

1. Consider the sailor-Boats-Reserve DB described as
 $s(sid, sname, rating, age)$, $b(bid, bname, color)$, $r(bid, sid, date)$
 Write each of the following queries in SQL and Relational algebra.

a) Find the ~~so~~ colours of boat reserved by Alber.

SELECT DISTINCT B.color
 FROM sailor S
 JOIN Reserve R ON S.sid = R.sid
 JOIN Boats B ON R.bid = B.bid
 WHERE S.sname = 'Alber';

$\pi_{color}(\sigma_{sname = 'Alber'}(S))$
 $\bowtie_{S.sid = R.sid}$
 $\bowtie_{R.bid = B.bid}$

b. Find all sailor ids of sailor who have a rating of at least 8 or reserved boat 103.

SELECT DISTINCT sid
 FROM sailor
 WHERE rating ≥ 8
 UNION

$\pi_{sid}(\sigma_{rating \geq 8}(S))$
 \cup
 $\pi_{sid}(\sigma_{bid = 103}(R))$

select Distinct sid
 From Reserve
 where bid = 103;

c. Find the names of sailor who have not reserved a boat whose name contains the string "storm". Order the names in ascending order.

select distinct s.sname
 from sailor S
 where S.sid Not in (
 (select R.sid
 from Reserve R
 Join Boats B on R.bid = B.bid
 where B.bname like '%storm%')
)
 Order by S.sname ASC;

StormBoat $\leftarrow \sigma_{bname \text{ like } '%storm\%'}(B)$
 StormRes \leftarrow StormBoat $\bowtie_{B.bid = R.bid} R$
 StormSailor $\leftarrow \pi_{sid}(\text{StormRes})$
 Allsailor $\leftarrow \pi_{sid, sname}(S)$
 NonStormSailor $\leftarrow \text{Allsailor} - (\text{StormSailor} \bowtie_{sid} sid)$
 Result $\leftarrow \pi_{sname}(\text{NonStormSailor})$

d. Find the sailor ids of sailor with age over 20 who have not reserved a boat whose name include the string "thunder".

Select distinct s.sid
 from Sailer s
 Where s.age > 20
 and s.sid not in (
 select R.sid
 from Reservu R
 Join Boats B on R.bid = B.bid
 where B.bname like '%thunder%');

Oldersailer $\leftarrow \sigma_{\text{age} > 20}(s)$
 Thunder Boat $\leftarrow \sigma_{\text{bname LIKE } '%thunder\%'}(B)$
 Thunder Res $\leftarrow \text{Thunder Boat} \bowtie B.\text{bid} = R.\text{bid } R$
 Thunder Sailer $\leftarrow \pi_{\text{sid}}(\text{Thunder Res})$

OlderNonThunders $\leftarrow \pi_{\text{sid}}(\text{Oldersailer}) - \text{Thunder Sailer}$

2. Define SELECT operation in Relational algebra.

The select operation (denoted by σ (Sigma))
 is used to select a subset of the tuple from a relation
 based on a selection condition.

* SE acts as a filter in the database queries,
 as it keeps only those tuple that satisfy the
 qualifying condition.

For example:

(a) select the Employee tuple whose department

number is 4:
 $\sigma_{\text{DNO}=4}(\text{EMPLOYEE})$

(b) select the employee tuple whose salary is greater
 than \$ 30,000:

$\sigma_{\text{salary} > 30,000}(\text{EMPLOYEE})$

* The selection condition is a Boolean (conditional)
 expression specified on the attribute of relation R.

* Tuples that make the condition true are selected
 and those make the condition false are filtered out.

* select operation produces a relation S that has the
 same schema as R.

* It is commutative in nature, hence a cascade of
 select operation may be applied.

* The number of tuple in the result of a SELECT
 is less than (or equal to) the number of tuple in the
 input relation R.

3. Discuss about trigger.

A trigger is a procedure that runs automati-
 cally when a certain event occurs in the DBMS.
 In many cases it is convenient to specify the type
 of action to be taken when certain events occur
 and when certain conditions are satisfied.

CREATE TRIGGER statement is used to implement such actions in SQL.

General form:

```
Create Trigger <name>
Before / After / <event>
For each row / For each statement
When (<condition>)
<action>
```

A trigger has three components.

i) Event: Trigger ~~happens~~ is activated when an event happens. (Insert, update, Delete)

Whether it is before the event or after the event.

ii) Condition (Optional): is activated if a mentioned condition is true.

iii) Action: These ~~are~~ ^{is} actions performed by Trigger.

Example: If the employee salary increased by more than 10%. Then increment the rank field by 1.

```
Create trigger Empsal
Before update of salary on Employee
for each row
Begin
```

```
If (:new.salary > (:old.salary * 1.1)) Then
    :new.rank := :old.rank + 1;
```

```
END IF;
```

```
END;
```

Where := is an assignment operator

11. Define UNION operation in Relational algebra?

Union (\cup) is a binary operation denoted by \cup . The result of $R \cup S$, is a relation that includes all the tuple that are either in R or in S or in both R and S , and Duplicate tuple are eliminated, providing distinct tuple.

The two operand Relation R and S must be "type compatible"

i.e., they must have same number of attributes or have same or compatible domains.

Example:

To retrieve the social security number of all employees who either work in department 5 ~~or~~ or directly supervise and employee who works in department 5.

$DEP5-EMPS \leftarrow \sigma_{DNO=5}(EMPLOYEES)$

$RESULT1 \leftarrow \pi_{SSN}(DEP5-EMPS)$

$RESULT2(SSN) \leftarrow \pi_{SUPERSSN}(DEP5-EMPS)$

$RESULT \leftarrow RESULT1 \cup RESULT2$

Result1
SSN
123456789
333445555
666884444
453453453

Result2
SSN
333445555
888665555

Result
SSN
123456789
333445555
666884444
453453453
888665555

- * $R_1(A_1, A_2, \dots, A_n)$ and $R_2(B_1, B_2, \dots, B_n)$ are type compatible if
- * they have the same number of attributes
- * The domains of corresponding attributes are type compatible
i.e., $\text{dom}(A_i) = \text{dom}(B_i)$ for $i=1, 2, \dots, n$

5. What is the use of group by clause?

The group by clause groups rows that have the same values in specified columns into summary rows.

It's most often used with aggregate functions, like:

count() - number of items

sum() - total of values

avg() - average

min()/max() - minimum/maximum

Syntax:

Select column1, AGG_FUNC(column2)

from table

Group by column1;

- * Every column in the select must either be

~~Example:~~

a part of the Group by, or
used in an aggregate function.

- * You can use having to filter groups after aggregation.

Example:

select department, Avg(salary) as avg-salary
from employees

Result:

two columns of attribute department and avg-salary, where the average of each salary is the tuples under each department.

6. List the aggregate functions supported by SQL? //

Explain about Aggregate operators in sql with example.

Aggregate functions are used to summarize information from multiple tuples into a single-tuple summary. A number of built-in aggregate functions exist:

count, sum, max, min and avg.

The count functions returns the number of tuples or values as specified in a query.

The functions sum, max, min and avg can be applied to a set or multiset of numeric values and return, respectively the sum, maximum value, minimum value and average (mean) of those values. These functions can be used in the select clause or in a having clause.

Example:

i> Find the sum of the salaries of all employees, the maximum salary, the minimum salary, and the average salary.


```

select sum (salary), max (salary), min (salary),
       avg (salary)
from employee;

```

7. Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```

select sum (salary), max (salary), min (salary),
       avg (salary)
from (Employee join Department on Dno = Dnumbu)
where Dname = 'Research';

```

7. Discuss the basic form of SQL query using group and having clause?

Having provides a condition on the summary information regarding the group of tuples associated with each value of the grouping attribute. Only the groups that satisfy the condition are retrieved in the result of the query.

Example:

For each project on which more than two employees work, retrieve the project number, the project name, and the number of employees who work on the project.

```

select pnumbu, Pname, count (*)
from Project, works_on
where Pnumbu = Pno
Group by Pnumbu, Pname
having count (*) > 2;

```

pnumbu and pname from project table/relation is selected and by the condition $Pnumbu = Pno$ (Project) (works_on)

A new resultant table consisting of Pnumbu and pname, which has numbu of employee count more than 2, is displayed.

8. Define CROSS PRODUCT operation in Relational algebra.

This operation is used to combine tuples from two relations in a combinatorial fashion.

It is denoted by $R(A_1, A_2, \dots, A_n) \times S(B_1, B_2, \dots, B_m)$

* Result obtained is a relation Q with degree n+m attributes: $Q(A_1, A_2, \dots, A_n, B_1, B_2, \dots, B_m)$, in that order.

* The resulting relation state has one tuple for each combination of tuples - one from R and one from S.

Hence, if R has n_R tuples (denoted as $|R| = n_R$), and S has n_S tuples, then $R \times S$ will have $n_R \times n_S$ tuples.

* The two operands do not have to be "type compatible".

* Cross product is not a meaningful operation.
But it can be made meaningful when followed by
other operation.

Example: for non-meaningful operation:-

Female_emps $\leftarrow \sigma_{sex='F'}(Employee)$

Empnames $\leftarrow \pi_{Fname, Lname, ssn}(Female_emps)$

Emp-dependents $\leftarrow Empnames \times Dependent$

\Rightarrow Here Emp-dependents will contain every combination

* To make it meaningful - in the above given situation,
to keep only combinations where the Dependent is
related to the Employee, we add a select operation as
follows:

Female_emps $\leftarrow \sigma_{sex='F'}(Employee)$

Empnames $\leftarrow \pi_{Fname, Lname, ssn}(Female_emps)$

Emp-dependents $\leftarrow Empnames \times Dependent$

Actual-deps $\leftarrow \sigma_{ssn=Essn}(Emp-dependents)$

result $\leftarrow \pi_{Fname, Lname, dependent_name}(Actual-deps)$

9. Discuss correlated Nested queries:

Some queries require that existing values in
the database be fetched and then used in a comparison
condition. Such queries can be conveniently formulated
by using nested queries, which are complex select from-
where blocks within the where clause of another query.
The other query is called the outer query.

Whenever a condition in where clause of a nested
query references some attribute of a relation declared in
the outer query, the two queries are said to be correlated.

Example:-

select E.Fname, E.Lname

From Employee As E

where E.ssn &N (select Essn

From Dependent as D

where E.Fname = D.Dependent-name

And E.sex = D.sex)

The nested query is evaluated once for each tuple
(or combination of tuples) in the ~~the~~ outer query. we
can think of query in above example as follows:
For each EMPLOYEE tuple, evaluate the nested query,
which retrieves the Essn values for all Dependent
tuples with the same sex and name as that Employee
tuple; if the ssn value of the Employee tuple is in
the result of the nested query, then select that
Employee tuple.

10. What are Views? How are they created give an example? What are the problems associated with view manipulation?

A view in SQL terminology is a single table that is derived from other tables. Other tables can be base tables or previously defined views.

* A view does not necessarily exist in physical form. It is considered to be a virtual table, in contrast to base tables, whose tuples are always physically stored in the database.

In SQL, the command to specify a view is Create View. The view is given a (virtual) table name (or view name), a list of attribute names, and a query to specify the contents of the view.

Example:

```
Create view work-on1
```

```
as select Fname, Lname, Pname, Houres
```

```
From Employee, Project, Work-on
```

```
Where SSN = E.SN and Pno = P.number
```

In this case, work-on1 inherits the name of the view attribute from the defining (base) tables Employee, Project and Work-on.

Problems associated with view manipulation:

* A view is supposed to be always up-to-date. If we modify the tuples in the base table on which the view is defined, the view must automatically reflect

these changes. Hence, the view is not realized or materialized at the time of view definition but rather at the time when we specify a query on the view. It is the responsibility of the DBMS and not the user to make sure that the view is kept up-to-date.

* One strategy, called query modification, involves modifying or transforming the view query into a query on the underlying base tables.

The disadvantage of this approach is that it is inefficient for views defined via complex queries that are time-consuming to execute, especially if multiple queries are going to be applied to the same view within a short period of time.

* If the view is not queried for a certain period of time, the system may then automatically remove the physical table (materialized view) and recompute it from scratch when future queries reference the view.

11. Write and explain a query to find the name of sailor who have reserved a red boat?

SQL Query:

Select distinct S.sname

From sailors S

Join Reserves R on S.sid = R.sid

Join Boats B on R.bid = B.bid

Where B.color = 'red';

* The sailor table is assigned to an alias S. This table contains details like sid, sname, rating, and age.

* Join Reserves joins the sailor and Reserves table using the sid field. It connects each sailor to their reservations.

* Then, the resultant is joined with Boats table using bid. Now we know which sailor reserved which

specific boat.

* And then those are filtered to include only reservations for boats that are red.

From the result, we select only the name of the sailor (sname), that even if a sailor reserved a red boat more than once, their name appears only once.

Example:

S.sname	B.colours
Alice	red
John	red
John	red

Final result \Rightarrow

S.sname
Alice
John

12. Explain set operations of Relational Algebra with examples.

SQL has directly incorporated some of the set operations from mathematical set theory, which are also part of relational algebra.

* set union (Union)

* set difference (Except) and

* set Intersection (Intersection)

The relations resulting from these set operations are sets of tuples, that is, duplicate tuples are eliminated from the result. These set operations apply only to union-compatible relations, so we must make sure that the two relations on which we apply the operation have the same attributes and that the attributes appear in the same order in both relations.

1. Union: \cup :- Result of $R \cup S$, is a relation that includes all tuples that are either in R or in S or in both R & S, and gives distinct elements/details.

Ex: $DepEmps \leftarrow \sigma_{dept=5}(Employee)$

$Result1 \leftarrow \pi_{ssn}(DepEmps)$

$Result2(ssn) \leftarrow \pi_{supssn}(DepEmps)$

$Result \leftarrow Result1 \cup Result2$

\Rightarrow To retrieve the social security number of all employees who either work in department 5 or directly supervise an employee who works in department 5.

i) Intersection (\cap) :- Result of the operation $R \cap S$ is a relation that includes all tuples that are in both R and S.

Example:

$\pi_{\text{country, city}}(\text{Customers})$

\cap

$\pi_{\text{country, city}}(\text{Branches})$

Project the country & city columns from both the Customers and Branches relations. Then return the common country and city tuples.

ii) Difference ($-$) :- Result of $R - S$ is a relation that includes all tuples that are in R but not in S.

Example:-

$\pi_{\text{EmployeeID, name}}(E)$

$-$
 $\pi_{\text{contractorID, name}}(C)$

Project name and EmployeeID from Employee relation and name and contractorID from contractor relation and ~~provide~~ returns,

EmployeeID, who are not in contractors relation.

13) Explain about the selection, Union, Projection, Rename, division and Cartesian product operations in relational algebra.

1) selection:

select operation (σ) is used to select a subset of the tuples from a relation based on a selection condition.

Example:

* select the Employee tuples whose department number is 4:

$\sigma_{\text{dno}=4}(\text{Employee})$

ii) Union (\cup) :

The result of $R \cup S$ is a relation that includes all tuples that are either in R or in S or in both R and S.

$\pi_{\text{EmployeeID, name}}(E)$

\cup

$\pi_{\text{EmployeeID, name}}(P)$

Employees who are working fulltime and parttime.

iii) Projection (π)

<attribute list> is the desired list of attributes from relation R. \rightarrow The project operation removes any duplicate tuples as the result of the project operation must be a set of tuples.

Example:

$\pi_{\text{name}, \text{Department}}(\text{Employee})$

Select the Name and Department column from the Employee relation

iv) Rename: $\rho(f)$

In some cases, we may want to rename the attributes of a relation or the relation name or both.

In such case, rename (ρ) is used in the queries.

Example:

to rename the EmployeeID to EmpID and Name to EmpName.

$\rho(\text{EmpID}, \text{EmpName}, \text{Department}, \text{salary})(\text{Employee})$

v) Division: (\div)

$R(Z) \div S(Z)$, where X subset Z . Let $Y = Z - X$

(and hence $Z = X \cup Y$); that is let Y be the set of attributes of R that are not attributes of S .

The result of Division is a relation $T(Y)$, that include a tuple t if tuple t_p appear in R with $t_p(Y)$, and with

$t_p(X) = t_s$ for every tuple t_s in S .

Ex:

$\pi_{\text{sid}}(R) \div B$

This gives sid values from R that are associated with every bid of B .

Cartesian product:

This operation is used to combine tuples from two relations in a combinational fashion.

$R(A_1, A_2, \dots, A_n) \times S(B_1, B_2, \dots, B_m)$

Example:

$\text{Employee} \times \text{Department}$

→ Every pair of one employee with one department.

14. Work and Explain a Query for finding the names of sailors who have reserved a Red or a Green Boat.

Query:

Select distinct s.sname
from sailor s

Join Reserves r on s.sid = r.sid

Join Boat b on r.bid = b.bid

where b.color IN ('Red', 'Green')

* Sailors table assigns to alias s. Then sailor and Reserves table is joined to find which sailor reserved which boat.

* The resultant is joined with Boat table, to get the details of the reserved boats, then it is filtered to only those boats that are either Red or Green.

And then distinct snames are chosen and then shown as result.

15. Write and explain a query for finding the color of Boat reserved by 'Lubber'.

Query:

```
SELECT DISTINCT b.color
FROM SAILOR S
JOIN RESERVE r ON s.sid = r.sid
JOIN BOAT b ON r.bid = b.bid
WHERE s.sname = 'Lubber';
```

- * SAILOR table is assigned to its alias S.
- * Then the sailon table is joined with Reserve table to find out which sailor reserved what boat.
- * The above resultant is joined with the Boats table to access details of the boat that were reserved.
- * Then it is filtered to only the sailor named 'Lubber'.
- * Then return with the distinct color of all boat that Lubber had reserved.

16. Discuss Normalization. Define First Normal Form. List Different Normal Forms.

Normalization is a technique of organizing the data in the database. It is a systematic approach of decomposing tables to eliminate data redundancy and undesirable characteristics like Insertion, update and Deletion Anomalies.

It is a multi step process that put data into tabular form by removing duplicated data from the relation table.

It can be considered as a "filtering" or "purification" process to make the design have successively better quality.

First normal form:

It is defined to disallow multivalued attributes, composite attributes, and their combination.

- * It states that the domain of an attribute must include only atomic (simple, indivisible) values and that the value of any attribute in a tuple must be a single value from the domain of that attribute.
- for example:

Emp-proj		Project	
Ssn	Ename	Pnum	Hours

This is nested relation, which is not allowed in first normal form. And can be resolved by

Creating separate relation called project with pnumbr and hour, along with project name.

Different

These ~~types of~~ Normal forms are.

i) First normal form \rightarrow composite / complex attribute

ii) Second normal form \rightarrow full functional dependency

iii) Third normal form \rightarrow transitive dependency

iv) Boyce-Codd Normal form - nontrivial functional dependency

v) ~~multivalued~~ ~~one~~

Fourth Normal form - Multivalued Functional dependency

vi) Fifth Normal form - Join Dependencies.

(14) Explain with the help of example of the following relational algebra operation. Join, Difference, Union, Select.

Join:

The sequence of Cartesian product followed by select is used quite commonly to identify and select related tuple from two relation.

The operation Join combine this sequence into a single operation.

General form of Join operation ~~is~~ on two relations $R(A_1, A_2, A_3, \dots, A_n)$ and $S(B_1, B_2, B_3, \dots, B_m)$ is:

$R \bowtie_{\text{Join condition}} S$

Example:- To get the manager's name from each department, we need to combine each Department tuple with the Employee tuple whose ssn value matches with mgrssn value in department tuple.
i.e.,

Dept_Mgr \leftarrow Department $\bowtie_{\text{mgrssn=ssn}}$ Employee.

18. Discuss the E.R to Relational mapping algorithm with example for each step.

Step 1:- Mapping of Regular Entity type.

- * Create a relation (table) for each regular entity.
- * Include all simple attribute.
- * Choose a primary key (if it's composite, use all parts)

Example:- For Employee, include attribute like ssn, name, Address etc. and make ssn the primary key.

Step 2:- Map weak Entity type.

- * Create a relation for each weak entity
- * Include all attribute plus the primary key of the owner as a foreign key.
- * Use the owner's key + weak entity's partial key as the composite primary key.

Ex: Dependent include, Dep_name, B-date, Essn (foreign key)

Primary key = {Essn, Dependent_name}

Step 3:- Map Binary 1:1 Relationships.

- * Foreign key Approach - Add the primary key of one entity to the other as a foreign key.

• Merged Relation - Merge both entities into one if both have total participation.

* (non-reference table) - create a separate relation (less common for 1:1)

Example:- Manager \rightarrow Add Mgr-ssn to department and Mgr-start date.

Step 4: Map Binary 1:N Relationships.

* Add the primary keys of the "1" side as a foreign key to the "N" side.

* Include any attribute of the relationship.

* Example: Work_for \rightarrow Add Dnumber as Dno to Employee.

Step 5: Map Binary M:N relations.

* Create a new relation, include the primary keys of both entities as foreign keys.

* Include any attribute of the relationship.

* The composite primary key is both foreign keys.

Example: Work_on \rightarrow has Essn, Pno and Hour.

Step 6: Map Multivalued Attributes.

* Create a separate relation for each multivalued attribute.

* Include the attribute and the primary key to the owning entity.

* Compaity. primary key = (owning entity's PK + multivalued attribute)

Ex: Locations \rightarrow create Dept-locations with Dnumber and Dlocation.

Step 7: Map N-ary Relationship Type.

* For relationships involving more than 2 entities.

\rightarrow Create a new relation, include the primary keys of all involved entities as foreign keys.

* Add any relationship attribute.

Example: Supply \rightarrow include Sname, Partno, proj-name.

11. Illustrate insert, delete, update, alter, grant, revoke and drop commands in SQL.

i) Insert:- The insert operation provides a list of attribute values for a new tuple that is to be inserted into a relation R.

This violates any of the four types of constraints:-

Domain constraint, key constraint, Entity integrity, Referential integrity.

Ex: Insert <('Cecilia', 'F', 'Kolonsky', NULL, '1960-04-05', '6357 Windy Lane, Katy, TX', 'F', 2800, NULL, 4)>

Result:- This insertion violates the entity integrity constraint (NULL for the primary key ssn), so it is rejected.

\rightarrow (Correct insertion)

Insert <('Cecilia', 'F', 'Kolonsky', '677678989', '1960-04-05', '6357 Windy Lane, Katy, TX', 'F', 28000, NULL, 4)>

ii) delete:- Delete operation is used to delete any tuple from the relation R.

To specify deletion, a condition on the attribute of the relation select the tuple (or tuples) to be deleted.

Example:-

Delete the worker-on tuple with Essn = '999887777' and Pno = 10.

This deletes exactly one tuple, whereas

Delete the worker-on tuple with Essn = '999887777'

will delete all the tuple with the Essn, no matter the ~~value~~ value of project number in it.

iii> update:- (modify) used to change the values of one or more attributes in a tuple (or tuples) of some relation R. It is necessary to specify a condition on the attributes of the relation to select the tuple to be modified.

Example:-

Update the salary of the Employee tuple with ssn = '999887777' to 28000.

Here, the tuple of Employee relation, where the attribute salary corresponding to the ssn given, will be updated to the value 28000.

iv> Alter:-

The definition of a base table or of other named schema elements can be changed by using the ~~alter~~ Alter command. For base table, the possible alter table actions include adding or dropping a column (attribute), changing a column definition, or adding or dropping table constraint.

Ex:-

Alter table company.Employee add column Job VARCHAR(10)

=> Another attribute Job is added to the relation Employee.

v> grant:-

The grant command is used in SQL to give permission to a user, role or application so that they can do specific actions on database objects like tables, views or stored procedure.

structure:-

Grant <privilege>

on <object>

To <user/role>

[WITH GRANT OPTION];

<privilege> - actions like select, insert, update, delete etc.

<object> - the table or other object the permission apply to

<user/role> - who you're granting permission to

vi> Revoke:-

Revoke command is used in SQL to remove permission that were previously granted to a user, role or group.

Revoke <privilege>

on <object>

From <user/role>

<privilege> - permissions you want to take away

<object> - Table or other object

<user/role> - user or role that currently has the privilege

vii) drop:- drop command can be used to drop named schema elements, such as tables, domains, or constraints.

If a whole schema is no longer needed:-

Drop schema company cascade / Restrict

Cascade:- remove all the elements in it too

Restrict:- only when the schema is empty.

To drop a table:-

Drop table Dependent cascade.

To drop a constraint:-

If one attribute already has a constraint,

Drop constraint pk-student

20. Explain informal design guideline for relational schema design.

We have four informal design guidelines for relation schemas.

i) Guideline 1:- Making sure that the semantic of the attribute is clear in the schema.

Informally, each tuple in a relation should

represent one entity or relationship instance.

* Attributes of different entities should not be mixed in the same relation.

* If it is required, it should be foreign keys used to refer to other entities.

ii) Guideline 2:- Reducing the redundant information in tuple.

Designing a schema that does not suffer from the insertion, deletion and update anomalies.

Ex:- Emp-proj (Emp#, Proj#, Enam, Pname, No-hours)

changing a single name of project no Pj will result in changing ^{that of} all the employees working on that project

iii) Cannot insert a project unless an employee is assigned to it, and cannot insert an employee unless he/she is assigned to a project.

ii) When one project is deleted, it will result in deletion of all the employees working on that project.

iii) Guideline 3: Reducing the Null value in tuple.

Relations should be designed such that their tuples will have as few Null values as possible.

Attributes that are Null frequently could be placed in separate relations.

Null values is seen when, attribute not applicable for a tuple. or when the value is unknown or if it is not available.

iv) Guideline 4: Disallowing the possibility of generating spurious tuples.

The relations should be designed to satisfy the lossy join condition i.e., No spurious tuples should be generated by doing a natural join of any two relations.

That means, meaningful results are to be generated instead of unwanted, space consuming, redundant values.

21. Discuss the Equijoin and Natural Join with suitable example.

Equijoin:-

The most common use of join involve join conditions with equality comparisons only. Such a join, where the only comparison operator used is '=', is called an equijoin.

Example:-

Select student.name, student.id, record.clau, record.city
From student, record
Join record
on student.city = record.city.

Natural join:-

The two join attributes, or each pair of corresponding join attributes, have the same name in both relations.

To apply a natural join on the Dnumber attributes of Department and Dept_location,

Dept_locs \leftarrow Department * Dept_location.

- * Only attribute with the same name is Dnumber
- * An implicit join condition is created based on this attribute: location.Dnumber = Dept_location.Dnumber

22. What is Normalization? What are the conditions that are required for a relation to be in 1NF, 2NF. (16)

Second normal form:- (2NF) is based on the concept of full functional dependency.

A functional dependency $X \rightarrow Y$ is a full functional dependency if removal of any attribute A from X means that the dependency does not hold any more; that is, for any attribute A \in X, $(X - \{A\})$ does not functionally determine Y.

- * the test for 2NF involves testing for functional dependencies whose left-hand side attributes are part of the primary key.

Definition:- A relation schema R is in 2NF if every nonprime attribute A in R is fully functionally dependent on the primary key of R.

Example:-

~~Hour is fully dependent on the ssn and Pnumber.~~
Hour is fully dependent of the ssn and Pnumber, as hours of working is assigned particularly for the ssn and the project the employee is working on.

23. What is meant by functional dependencies?

A functional Dependency is a constraint between two sets of attribute from the database.

Given a relation R , a set of attribute X in R is said to functionally determine another attribute Y , also in R , if and only if each X value is associated with at most one Y value.

Consider the relation schema Emp-proj.

Emp-proj.

Ssn	Pnumber	Hours	Ename	Pname	Plocation
-----	---------	-------	-------	-------	-----------

function dependencies in there are:


- * $Ssn \rightarrow Ename$
- * $Pnumber \rightarrow \{Pname, Plocation\}$
- * $\{Ssn, Pnumber\} \rightarrow Hours$

* The value of an employee's social security number (Ssn) uniquely determines the employee name (Ename).

* The value of a project's number (Pnumber) uniquely determines the project name (Pname) and location (Plocation).

* A combination of Ssn and Pnumber values uniquely determines the number of hours the employee currently work on the project.

Emp-proj



24. Explain the type of anomalies with example.

Anomalies are classified into:

- i) Insertion anomaly
- ii) deletion anomaly
- iii) modification anomaly.

i) Insertion anomaly:-

Difficulty in inserting data because some other data is also required.

Example:- In an Emp-Dept (which store both employee and department info), if we have to add a new department, that doesn't yet have any employee, it is an insertion anomaly because, the primary key-ssn in employee table cannot be NULL, violating entity integrity.

ii) Deletion anomaly:-

Losing useful data because another piece of data was deleted.

Example:- If we delete the last employee working in a department from Emp-dept, we lose all information about the department (like its name, manager, location) since it's stored redundantly with each employee.

* The system becomes fragile and prone to data loss.

iii) Modification Anomaly:

~~We~~ Need to change the same data in multiple places; Failure to do so leads to inconsistency.

Example:- If the manager of a department changes, we have to update that manager's ID in every row of the Emp-dept table for employees in that department.

If it fails, we have inconsistent data which leads to error and ^{more} time consumption.

25. Define trigger, and explain its three parts, compare row level and statement level trigger. (3)

Row Level	Statement level
Row level trigger executes once for each and every row in the transaction.	Statement level trigger executes only once for each single transaction.
Specifically used for data auditing purpose.	Used for enforcing all additional security on the transactions performed on the table.
"FOR EACH ROW" clause is present in <u>Create trigger</u> command.	"FOR EACH STATEMENT" clause is is omitted in <u>Create trigger</u> command.
Ex: If 1500 rows are to be inserted into a table, the row level trigger would execute 1500 times.	Ex: If 1500 rows are to be inserted into a table, the statement level trigger would execute only once.

26. Demonstrate transitive dependency? Give an example.

A functional dependency $X \rightarrow Y$ in a relation schema R is a transitive dependency if there exist a set of attributes Z that are neither a primary nor a subset of any key of R (candidate key) and both $X \rightarrow Z$ and $Z \rightarrow Y$ hold.

Emp_dept:-

Ename	Ssn	Bdate	Address	Dnumber	Dname	Dmgrssn
	↑	↑	↑	↑	↑	↑

$ssn \rightarrow Dmgrssn$ is a transitive functional dependency since $ssn \rightarrow Dnumber$ and $Dnumber \rightarrow Dmgrssn$ hold.

$Dnumber$ is neither a key itself nor a subset of the key of EMP_Dept.

* ~~As a relation schema is in 3NF~~

* $ssn \rightarrow Ename$ is non-transitive since there is no set of attributes X where $ssn \rightarrow X$ and $X \rightarrow Ename$.

27. Explain BCNF. What are the steps to be followed to convert a relation in 3NF to BCNF?

Boyce-Codd Normal Form, was proposed as a simple form of 3NF. Every relation in BCNF is also in 3NF; however, a relation in 3NF is not necessarily in BCNF.

⇒ A relation schema R is in BCNF if whenever a nontrivial functional dependency $X \rightarrow A$ holds in R , then X is a superkey of R .

Only if $X \rightarrow A$ holds in a relation schema R with X not being a superkey and A being a prime attribute will R be in 3NF but not in BCNF.

Example:-

Consider a relation

Course (CourseID, Instructor, Room)

~~Here~~, functional dependencies hold:-

CourseID \rightarrow Instructor

Room \rightarrow Instructor

Here, CourseID is the primary key, so the relation is in 3NF, but not in BCNF.

Because $Room \rightarrow Instructor$ but Room is not a superkey ~~primary key~~ but it determines another attribute (Instructor), which violates BCNF.

To make this a BCNF, we decompose course relation into two relations.

is Room_Instructor (Room, Instructor) which handles the dependency $Room \rightarrow Instructor$.

17. Course-Info: (CourseID, Room) which keeps the link between CourseID and Room.

⇒ Room-Instructor

Room	Instructor	CourseID	Room
101	Alice	CS101	101
102	Bob	CS102	102

This has \times No Redundant data

* No anomalies

* Both relations are in BCNF.

28. What is normalization? What are the conditions that are required for a relation to be in 1NF, 2NF & 3NF. explain with example.

3NF:-

A relation schema R is in third normal form (3NF), if, whenever a nontrivial functional dependency $X \rightarrow A$ holds in R , either (a) X is superkey of R , or (b) A is a prime attribute of R .

consider a relation LOTS1

	Property-Id	Country-name	Lot#	Area	Price
FD1					
FD2					
FD3					

This is not in 3NF, because Area is not a superkey and Price is not a prime attribute in LOTS1

To normalize it into 3NF, we decompose it into the relation schemas LOTS1A and LOTS1B

LOTS1A

	PropertyId#	Countrysname	Lot#	Area
FD1				
FD2				

LOTS1B

	Area	Price
FD3		

* Now both LOTS1A and LOTS1B are in 3NF.